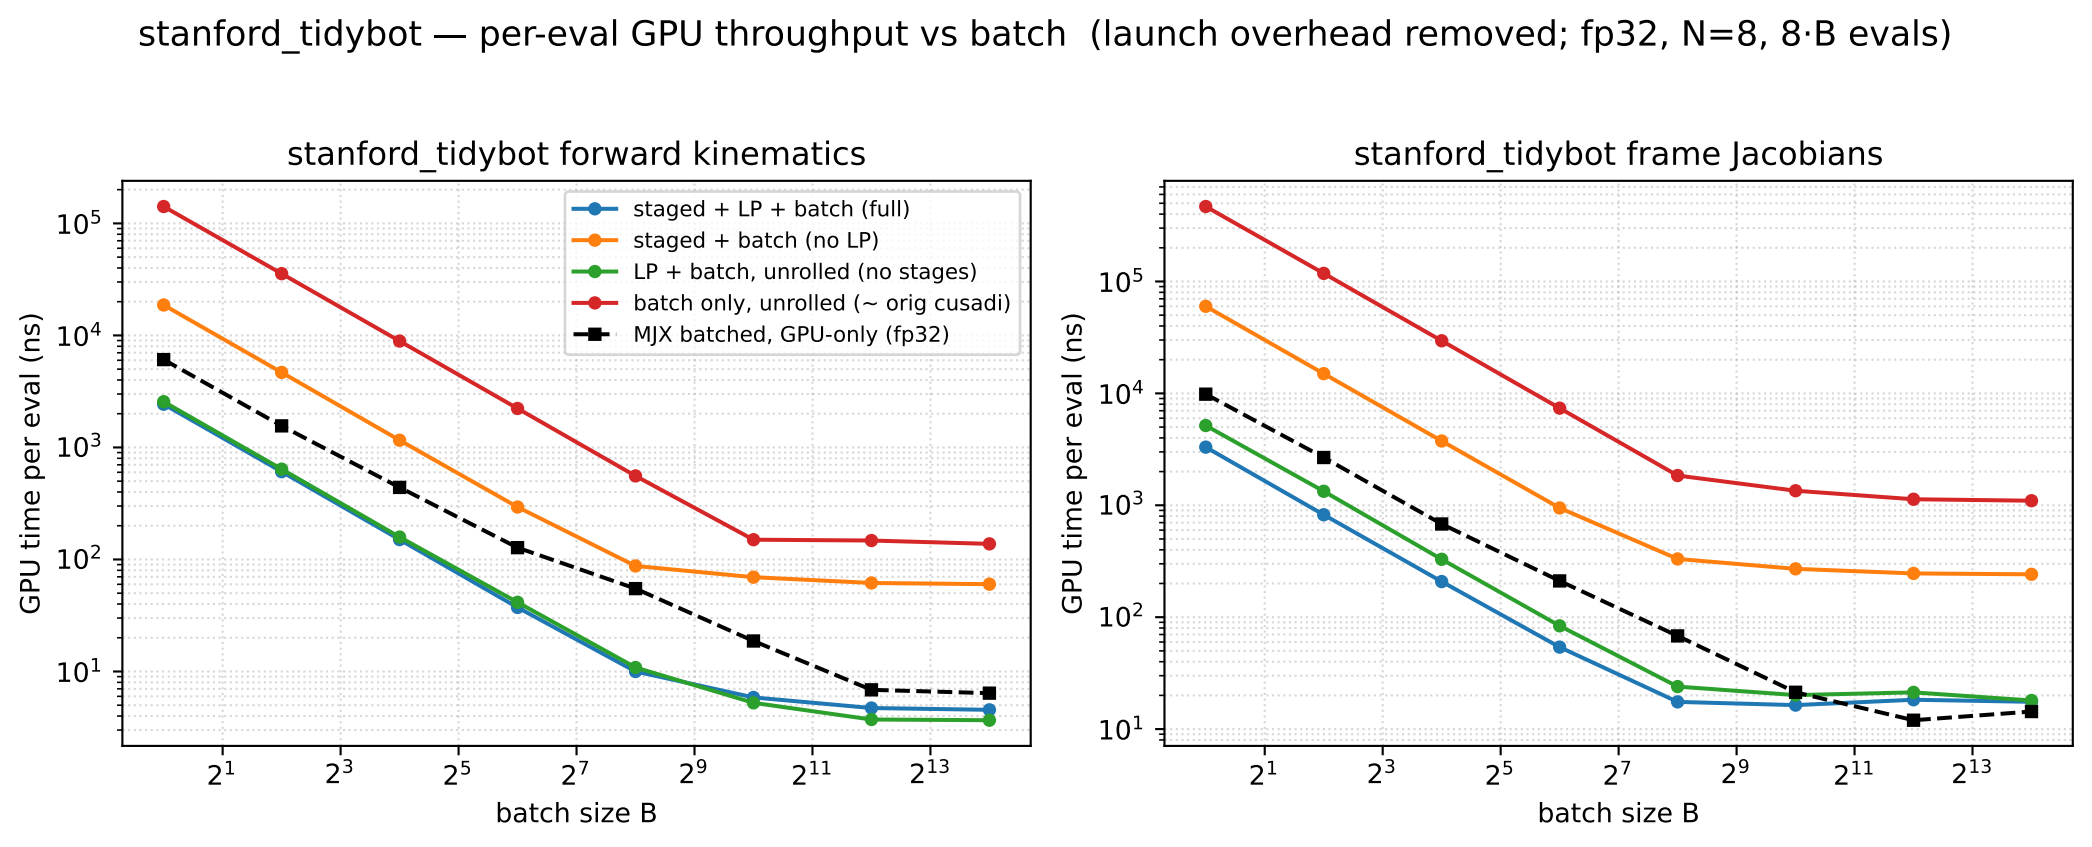
stanford_tidybot — per-eval GPU throughput vs batch (launch overhead removed; fp32, N=8, 8·B evals)
stanford_tidybot forward kinematics
stanford_tidybot frame Jacobians
staged + LP + batch (full)
staged + batch (no LP)
LP + batch, unrolled (no stages)
batch only, unrolled (~ orig cusadi)
MJX batched, GPU-only (fp32)
GPU time per eval (ns)
batch size B